

High Diminsional Statistics-Sheet 4-Exercise 4

Hamed Vaheb

02/05/2022

Use the same prostate dataset as in the last sheet, which can be found at [this clickable link](https://hastie.su.domains/ElemStatLearn/datasets/prostate.data)

We want to perform the regularization method known as Adaptive Lasso, which seeks to minimize

$$RSS(b) + \lambda \sum_{j=1}^p \hat{w}_j |b_j|,$$

where $RSS(b)$ is the residual sum of square of b , λ is the tuning parameter (chosen through 10-fold cross validation), b_j are the p estimated coefficients and \hat{w}_j are adaptive weights. We recall that

$$\hat{w}_j = \frac{1}{(|b_j^{in}|)^\gamma},$$

where b_j^{in} is an initial estimate of the coefficients and γ is a positive constant for adjustment of the adaptive weights.

Prepare

First of all we import the dataset:

```
url <- "https://hastie.su.domains/ElemStatLearn/datasets/prostate.data"
df <- read.table(url, sep = '\t', header = TRUE)
df %>% head(10)
```

```
##      X      lcavol  lweight age      lbph svi      lcp gleason pgg45      lpsa
## 1    1 -0.5798185 2.769459  50 -1.3862944  0 -1.386294      6      0 -0.4307829
## 2    2 -0.9942523 3.319626  58 -1.3862944  0 -1.386294      6      0 -0.1625189
## 3    3 -0.5108256 2.691243  74 -1.3862944  0 -1.386294      7     20 -0.1625189
## 4    4 -1.2039728 3.282789  58 -1.3862944  0 -1.386294      6      0 -0.1625189
## 5    5  0.7514161 3.432373  62 -1.3862944  0 -1.386294      6      0  0.3715636
## 6    6 -1.0498221 3.228826  50 -1.3862944  0 -1.386294      6      0  0.7654678
## 7    7  0.7371641 3.473518  64  0.6151856  0 -1.386294      6      0  0.7654678
## 8    8  0.6931472 3.539509  58  1.5368672  0 -1.386294      6      0  0.8544153
## 9    9 -0.7765288 3.539509  47 -1.3862944  0 -1.386294      6      0  1.0473190
## 10  10  0.2231436 3.244544  63 -1.3862944  0 -1.386294      6      0  1.0473190
##      train
## 1      TRUE
## 2      TRUE
## 3      TRUE
## 4      TRUE
## 5      TRUE
## 6      TRUE
## 7     FALSE
## 8      TRUE
## 9     FALSE
## 10     FALSE
```

Question 1

Build the regression model for the variable prostate antigen (lpsa). Use it to obtain the initial estimates of the coefficients b_j^{in} .

At first we build the regression model:

```
##
## Call:
## lm(formula = df$lpsa ~ ., data = features)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.76644 -0.35510 -0.00328  0.38087  1.55770
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.181561   1.320568   0.137  0.89096
## lcavol       0.564341   0.087833   6.425 6.55e-09 ***
## lweight      0.622020   0.200897   3.096  0.00263 **
## age         -0.021248   0.011084  -1.917  0.05848 .
## lbph         0.096713   0.057913   1.670  0.09848 .
## svi          0.761673   0.241176   3.158  0.00218 **
## lcp         -0.106051   0.089868  -1.180  0.24115
## gleason      0.049228   0.155341   0.317  0.75207
## pgg45        0.004458   0.004365   1.021  0.31000
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6995 on 88 degrees of freedom
## Multiple R-squared:  0.6634, Adjusted R-squared:  0.6328
## F-statistic: 21.68 on 8 and 88 DF,  p-value: < 2.2e-16
```

The estimators of b_0 and $b_j \in \{1, \dots, 8\}$ are respectively:

```
coef(model_lm)

## (Intercept)      lcavol      lweight      age      lbph      svi
## 0.181560845  0.564341280  0.622019788 -0.021248185  0.096712522  0.761673402
##          lcp      gleason      pgg45
## -0.106050939  0.049227934  0.004457512
```

Perform lasso regression with 10-fold cross validation to find the best λ , from which we will create the initial estimates the coefficients

```
X <- data.matrix(features)
y <- df$lpsa

cv_lasso <- cv.glmnet(X, y,
                      ## type.measure: loss to use for cross-validation.
                      ## K = 10 is the default.
                      nfold = 10,
                      ## 'alpha = 1' is the lasso penalty, and 'alpha = 0' the ridge penalty.
                      alpha = 1)
## Penalty vs CV MSE plot
```

Store the coefficients of the lasso model in which we set λ to be `lambda.min` which is λ at which the smallest mean squared error (MSE) is achieved.

```
## s: Value(s) of the penalty parameter 'lambda' at which
## predictions are required. Default is the entire sequence used
## to create the model.
lasso_coef <- coef(cv_lasso, s = cv_lasso$lambda.min)
```

Question 2

Create the Adaptive Weights \hat{w}_j for $\gamma = 0.5, 1$ and 2 . Choose the best λ through 10-fold cross validation.

The procedure I pursue is as following:

1. Using initial estimates of b_j^{in} derived in Question 1, I Construct w (adaptive weights) by adjusting the `lasso_coef` using 3 three different choices for γ . For this purpose, I define a function `weight_func` which constructs adaptive weights given γ and b , and I make sure it doesn't contain ∞ as a value.
2. We define the function `alasso_cv_gamma` which takes γ as its only argument and builds a 10-fold cross-validation adaptive lasso model using adaptive weights that are created based on γ .
3. We compare the three models defines in step 2 in order to find the γ that leads to the "best" model, i.e., the model that has the λ that leads to the least mse loss.
4. As the ultimate goal is to find the best λ , I create my final cross-validation model with the best γ , and output the best λ . Moreover, I report the coefficients' estimates.
5. For the purpose of illustration, I plot values of coefficients for two model variations, i.e., simple lasso and adaptive lasso. Finally, I plot the lambda choices of the cross-validation model.

Step 1: Construct Adaptive Weights

```
#best_lasso_coef <- as.numeric(coef(cv_lasso, s = #cv_lasso$lambda.min))[-1]

## constructing w which consists of adaptive weights
## The intercept estimate should be dropped.
p <- nrow(lasso_coef)-1
b <- lasso_coef[1:p]
weight_func <- function(x,gamma){
  1/(abs(x)**gamma)
}
gamma <- c(1/2,1,2)
adap_weights <- sapply(b, weight_func,gamma[2])
adap_weights

## [1] 5.449856 1.819027 1.642293 53.076800 11.041996 1.389725 13.067723
## [8] 24.777622

# Replacing values estimated as Infinite for 999999999
adap_weights[adap_weights == Inf] <- 999999999
```

Step 2: Define the function `alasso_cv_gamma`

```
alasso_cv_gamma <- function(gamma) {
  adap_weights <- sapply(b, weight_func, gamma)
  adap_weights[adap_weights == Inf] <- 999999999
  alasso_cv <- cv.glmnet(X,y,nfold = 10,alpha = 1,penalty.factor = adap_weights,keep = TRUE)
  y_pred <- predict(alasso_cv, newx = X, s = alasso_cv$lambda.min)
  return(y_pred)
}
```

Step 3: Find the best gamma

```
## create a dictionary with keys as gamma and values as MSE
dict = c()
keys = c()
for (g in gamma)
{
  keys <- append(keys, as.numeric(g))
  y_pred <- alasso_cv_gamma(g)
  mse <- mean((y_pred - y)^2)
  dict[sprintf("%f",as.numeric(g))] <- as.numeric(mse)
}

gamma_best <- keys[which.min(dict)]
dict
```

```
## 0.500000 1.000000 2.000000
## 0.4439530 0.4714817 0.4656138
```

Step 4: Report the best lambda and coefficients' estimates

```
adap_weights <- sapply(b, weight_func, gamma_best)
adap_weights[adap_weights == Inf] <- 999999999
lasso_cv_best <- cv.glmnet(X,y, nfold = 10, alpha = 1, penalty.factor = adap_weights, keep = TRUE)
y_pred <- predict(lasso_cv_best, newx = X, s = lasso_cv_best$lambda.min)

#obtain the minimum lambda
lasso_cv_best$lambda.min


## [1] 0.004823694

#obtain the corresponding coefficients
coef(lasso_cv_best, s = lasso_cv_best$lambda.min)
```


```
## 9 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept)  0.096575102
## lcavol      0.556091641
## lweight     0.629922897
## age        -0.019741987
## lbph        0.085687021
## svi         0.731214190
## lcp        -0.088494073
## gleason     0.048489142
## pgg45       0.003927875
```

Step 5: Plotting the estimations of the coefficients for all models


```
model_lasso <- glmnet(features, df$lpsa, alpha = 1)
lasso <- glmnet(features, df$lpsa, alpha = 1, penalty.factor = adap_weights)
plot(model_lasso)
```



```
plot(lasso)
```



```
## Penalty vs CV MSE plot
plot(lasso_cv_best)
```



Question 3:

Use the `glmnet` function to execute the adaptive Lasso. Plot the area under ROC (receiver operating characteristic) curve, also called AUC, and report the values of minimum λ (obtained for minimum AUC).

We discretize the target variable (ibsa), i.e., convert it from continuous type to discrete one so as to use the metric AUC (since it is a classification metric). We do the same for the predictions. Then, we can compare them using AUC metric. Finally, we plot the ROC-AUC curve.

Step 1: Discretize the target variable (ibsa)

```
y_disc <- as.integer(discretize(y, breaks = 2, labels=c(0, 1)))
y_pred_disc <- as.integer(discretize(y_pred, breaks = 2, labels=c(0, 1)))
```

Step 2: Calculate AUC

```
## Extract predicted probabilities and observed outcomes.
## pROC for ROC construction
roc <- pROC::roc(y_disc ~ y_pred_disc)

## Setting levels: control = 1, case = 2
## Setting direction: controls < cases
auc <- auc(y_disc, y_pred)

## Setting levels: control = 1, case = 2
## Warning in roc.default(response, predictor, auc = TRUE, ...): Deprecated use a
## matrix as predictor. Unexpected results may be produced, please pass a numeric
## vector.
## Setting direction: controls < cases
```

```
auc
```

```
## Area under the curve: 0.8971
```

Step 3: Plot ROC-AUC Curve

```
## Plot an ROC curve with AUC and threshold
```

```
plot(roc, print.auc = TRUE, print.thres = TRUE, print.thres.best.method = "youden")
```

